

Product

A Product for Sub-Quadratic Attention Research

Whitepaper — Draft v0.1

July 2026

Abstract

Product is a proposed commercial venture whose mission is to discover, validate, and continuously improve efficient replacements for the quadratic self-attention mechanism that bottlenecks all modern Large Language Models (LLMs). Contributors compete by submitting open-source GPU kernels that reproduce the mathematical output of standard $O(n^2)$ attention while operating in sub-quadratic time and memory. Under a winner-take-all prize structure, the highest-scoring kernel each cycle captures the full prize pool for that round, driving intense, continuous optimization. An automated evaluation engine scores submissions through fully deterministic, hardware-measured benchmarks — no language-model judge is involved. The Product's commercial strategy centers on API monetization: enterprises pay to use the champion kernel in production, and that revenue is recycled back into the prize pool to deepen the competition and accelerate the Product's roadmap.

Table of Contents

1. The Problem — Why LLMs Hit a Wall
 2. Mission Statement
 3. Technical Background
 4. Product Architecture
 5. What Contributors Do
 6. How the Evaluation Engine Works
 7. Anti-Gaming Mechanisms
 8. Incremental Difficulty Roadmap
 9. The Product Offering
 10. Business Model — Revenue & the Reinvestment Flywheel
 11. Why This Succeeds Where Others Have Failed
 12. Conclusion
-

1. The Problem — Why LLMs Hit a Wall

Modern Large Language Models are built on the Transformer architecture, which uses Self-Attention to let every word in a sequence look at every other word. This enables deep contextual understanding — but at a steep computational price.

The core self-attention operation requires computing a score matrix of size $n \times n$, where n is the number of input tokens. This means:

- At 8,000 tokens: ~64 million score values per attention layer
- At 128,000 tokens: ~16 **billion** score values per layer
- At 1,000,000 tokens: ~1 **trillion** score values per layer — requiring hardware far beyond what is commercially available today

While recent advances such as extended context windows in frontier models have pushed the practical limit higher, they do so primarily through engineering workarounds — chunked attention, sparse approximations, and massive memory bandwidth — not by changing the underlying $O(n^2)$ math. This means cost and latency continue to grow non-linearly with context, making truly long-horizon AI agents and document understanding economically prohibitive at scale.

Meanwhile, alternative architectures that achieve $O(n)$ or $O(n \log n)$ complexity exist in research, but they trade accuracy for speed — they tend to forget details from earlier in the sequence, making them unsuitable for precision-critical enterprise applications.

The gap between quadratic accuracy and linear speed is the unsolved problem that Product is designed to close — through open, incentivized competition.

2. Mission Statement

Product's mission is to harness a global pool of compute and engineering talent to produce the world's most accurate, fastest, and most memory-efficient sub-quadratic attention kernels for Transformer-based LLMs — and to commercialize those kernels as a product that generates real revenue and funds its own continued improvement.

- **Replace $O(n^2)$ attention math** with sub-quadratic alternatives that preserve near-identical output accuracy.
- **Make the solution hardware-real** — contributors must write optimized GPU kernels (Triton/CUDA) that actually run faster on production hardware.

- **Keep all kernels open-source** — every submission is public, allowing the community to study, fork, and improve the current champion.
- **Apply winner-take-all pressure** — the highest-scoring kernel each cycle captures the full prize pool, forcing genuine best-effort competition.
- **Generate revenue that compounds** — enterprise API sales fund a growing prize pool, which attracts stronger contributors and produces better kernels.

3. Technical Background

3.1 How Standard Attention Works

In a Transformer model, the input token embeddings (X) are projected through three learned weight matrices into Query (Q), Key (K), and Value (V):

- **Query ($Q = X \cdot W_Q$)**: "What am I looking for?" — the search intent of each token.
- **Key ($K = X \cdot W_K$)**: "What do I offer?" — the label each token presents to others.
- **Value ($V = X \cdot W_V$)**: "What is my content?" — the actual information each token shares.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{d}}\right) \cdot V$$

The critical step is $Q \cdot K^T$, which produces the $n \times n$ score matrix. This single operation is where both memory and compute explode quadratically. The final output has the same shape as the input $[n, d]$ — the bottleneck exists only in the intermediate step, which means a drop-in replacement is mathematically possible.

3.2 The KV Cache Compounds the Problem

During autoregressive inference, models cache previously computed Keys and Values (the KV Cache) to avoid re-processing the entire history at each generation step. As conversations or documents grow, this cache can exceed the model's weight size, saturating GPU VRAM and degrading throughput significantly — a second-order effect of the same quadratic root cause.

3.3 The Sub-Quadratic Opportunity

Because the input and output shapes of the attention layer are identical $([n, d])$, any algorithm that maps $Q, K, V \rightarrow \text{output}$ without materializing the $n \times n$ matrix is a valid drop-in replacement. Approaches such as spectral methods, hierarchical summaries, and FFT-based mixing have demonstrated this is feasible in research — but none have been refined into production-grade kernels that match Transformer accuracy at scale. That is the gap this Product fills.

4. Product Architecture

Product is built as a continuous, open engineering competition wrapped around a commercial API. The evaluation logic — the specification defining the contributor and scoring interfaces — is maintained and operated by the Product team. Each evaluation cycle runs for a fixed interval of approximately 72 minutes. At the end of each cycle, an automated settlement system aggregates benchmark results and awards the entire prize pool to the top-ranked contributor.

The size of the prize pool for each cycle is driven directly by product revenue: the more enterprises adopt the champion kernel, the larger the pool, and the stronger the incentive for contributors. This makes commercial success a direct lever on the rate of technical progress — revenue and capability reinforce each other in a single loop.

Component	Role	Key Output
Evaluation Engine	Issues challenges, runs benchmarks, ranks contributors, and settles the prize pool each cycle	Contributor scores → prize payout
Contributor	Develops and submits open-source sub-quadratic attention kernels each cycle	Kernel source code + result tensors
Champion Archive	Stores the top-ranked open-source kernel each round	Public GitHub library
Product Team	Maintains the evaluation specification; runs the commercial API; reinvests revenue into the prize pool	Operating margin + API revenue

The evaluation mechanism is fully deterministic. Scoring uses no LLM judges — ranking is based on MSE error, VRAM usage, and wall-clock latency measured on standardized hardware.

5. What Contributors Do

5.1 The Core Task

Each contributor is a GPU engineer and researcher. Their goal is to submit an open-source attention kernel — code that accepts the standard Q, K, V input tensors from an existing model and returns the correct contextualized output tensor, but does so without constructing the full $n \times n$ score matrix.

5.2 Open-Source Requirement

All kernel submissions must be published as open-source code (MIT or Apache 2.0 license) in the Product's public repository before the cycle closes. This requirement serves two purposes: it allows the evaluation engine to independently re-compile and re-run the code as the ground truth for performance claims, and it enables every other contributor to study the current champion and attempt to improve upon it. Closed or obfuscated submissions are automatically disqualified.

5.3 The Input Contributors Receive

- **Q, K, V tensors** of shape $[n, d]$, derived from a reference model's weights. Tensors are randomly seeded each round — contributors cannot cache or precompute answers.
- **Sequence length target (n)** — increases over time as the benchmark difficulty ramps up across phases.
- **Precision budget** — the maximum allowed Mean Squared Error (MSE) between the contributor's output and the ground-truth output.

5.4 The Output Contributors Submit

- **Result tensor**: the computed attention output of shape $[n, d]$.
- **Open-source kernel code**: Triton, CUDA, or equivalent. Published to the Product repository; independently re-compiled by the evaluation engine to verify performance claims.
- **Reproducibility metadata**: hardware spec, compilation flags, runtime configuration.

5.5 Winner-Take-All Incentive Structure

To avoid score ceilings that let contributors stagnate at a merely satisfactory level, Product enforces a winner-take-all structure: the contributor whose kernel achieves the highest composite efficiency score in a given cycle receives 100% of the prize pool for that cycle. All other contributors receive zero.

This creates maximum competitive pressure. Because the champion kernel is open-source, losing contributors can study exactly what beat them and submit an improved version the next cycle. The Product therefore converges via public, evolutionary competition rather than private, parallel guessing.

Winner-take-all + open-source = a ratchet mechanism. Each champion raises the bar for everyone, and the next winner must genuinely surpass it.

6. How the Evaluation Engine Works

The evaluation engine runs the scoring pipeline independently at the end of each cycle. Scoring is fully automated and deterministic, so every contributor can reproduce and audit their own result.

6.1 Tier 1 — Accuracy Gate (Pass/Fail)

The engine computes the ground-truth output using the standard $O(n^2)$ formula on the same Q, K, V tensors at a tractable small n . It then measures the Mean Squared Error (MSE) between the ground truth and the contributor's returned tensor.

If $\text{MSE} > \text{threshold} \rightarrow \text{score} = 0$. The kernel is producing inaccurate results and earns no reward for this cycle.

6.2 Tier 2 — Scaling Proof

The engine independently re-runs the contributor's submitted kernel on standardized hardware at a large sequence length (e.g., $n = 128\text{K}$ or $n = 256\text{K}$):

- **Peak VRAM measurement:** a true $O(n^2)$ kernel will exhaust memory or use exponentially more VRAM than a sub-quadratic one.
- **Wall-clock latency:** measured across multiple runs to eliminate noise.
- **Scaling curve check:** the kernel is also run at $2\times$ the sequence length. If latency quadruples, the contributor has not achieved sub-quadratic scaling.

6.3 Tier 3 — Composite Score & Winner Selection

Among contributors who pass both the accuracy and scaling gates, the composite score is:

$$\text{Score} = (1 / \text{MSE_error}) \times (1 / \text{Peak_VRAM}) \times (1 / \text{Latency})$$

The contributor with the highest score wins the full prize pool for the cycle. All others receive nothing. The engine records the result and publishes the champion kernel to the archive.

6.4 Periodic Full-Model Spot Check

Periodically, the engine performs a full-model spot check: a unique fact is injected deep in a long sequence, and the model — running the contributor's kernel across all layers — must retrieve it. This verifies that approximation errors do not compound across the model's 32–80 layers to the point of information loss.

7. Anti-Gaming Mechanisms

Any competition with real money on the line will attract contributors who optimize for the score rather than the goal. Product's design closes all major gaming vectors:

Attack Vector	Defense Mechanism
Submitting standard $O(n^2)$ code	Scaling curve test: latency at $2\times$ sequence length reveals quadratic growth immediately.
Hardcoding answers for known inputs	Q, K, V tensors are randomly seeded every cycle — no two rounds share the same inputs.
Claiming performance without running the code	The evaluation engine independently re-compiles and re-runs submitted open-source code on its own hardware.
Copying the champion kernel verbatim	Functional fingerprinting: identical computational graphs are flagged. The open-source requirement also means copying is visible — contributors must demonstrably improve the code.
Using a faster personal GPU to appear efficient	All measurements are taken by the engine on locked, standardized hardware environments.
Submitting a lossy approximation	The MSE gate at Tier 1 is a hard filter. Kernels that deviate beyond the precision budget earn zero regardless of speed.

8. Incremental Difficulty Roadmap

The benchmark difficulty increases as contributors prove each phase is solved. This keeps short-term rewards accessible to skilled engineers while ensuring the Product never stagnates. Each phase transition is triggered by achievement, not by calendar.

Phase	Context Length	MSE Threshold	Contributor Profile	Timeline
1 — Foundation	Up to 32K tokens	< 0.01	Adapt existing open-source optimizations (FlashAttention variants, sparse attention) to the Product format	Months 1–2
2 — Log-Linear	Up to 128K tokens	< 0.005	Novel hybrid approaches required; basic spectral or low-rank kernels needed to pass the scaling gate	Months 3–5
3 — Million-Token	Up to 1M tokens	< 0.001	Production-grade kernels with true sub-quadratic VRAM scaling; hardware-specific Triton/CUDA tuning	Months 6–12
4 — Convergence	Arbitrary	< 0.0005	Fine-grained hardware tuning; optional weight co-optimization (LoRA patches) to recover residual approximation error	Ongoing

9. The Product Offering

9.1 The Open-Source Kernel Library

The primary output of Product is a public, open-source library of sub-quadratic attention kernels. At the end of each cycle, the champion kernel is published to the Product's GitHub repository — the fastest, most accurate, most memory-efficient implementation yet validated by hardware benchmark.

Any developer running Llama-3, Mistral, Qwen, or any other Transformer-based model can swap their attention layer for a Product champion kernel and gain:

- Longer effective context windows without adding hardware
- Higher inference throughput (tokens per second) at the same GPU budget
- Reduced per-token inference cost for long-document and agentic workloads

9.2 The Commercial API

The Product team operates a commercial API that wraps the current champion kernel. Enterprise customers submit their model configuration and receive an optimized, tested attention module in return. The API handles compatibility testing, regression checks, and versioning — removing the integration burden from the customer.

- AI inference providers seeking lower latency at long context (vLLM, SGLang, Fireworks, Together.ai)
- Enterprise teams processing long legal, medical, financial, or scientific documents
- AI agent infrastructure companies requiring sustained multi-turn memory without context degradation

10. Business Model — Revenue & the Reinvestment Flywheel

The core commercial thesis is simple: generate real revenue from enterprises that need efficient LLM inference, and recycle a portion of that revenue back into the prize pool. This creates a direct feedback loop between product success and the pace of technical progress.

10.1 The Revenue Engine

Product charges enterprises for access to the validated champion kernel via a tiered API subscription:

Tier	Offering	Target Customer
Standard	Monthly API access to current champion kernel + integration docs	Startups, mid-size AI teams
Pro	Standard + priority updates, SLA guarantees, compatibility testing for specific model families	Inference providers, enterprise teams
Enterprise	Custom integration, private deployment, co-development of model-specific weight patches	AI labs, cloud providers, defense

The value proposition is strong: a 10–20% reduction in GPU inference costs at long context represents millions of dollars annually for any organization running LLMs at scale. API pricing can therefore be set well above the marginal cost of delivery while remaining compelling to customers.

10.2 The Reinvestment Flywheel

A defined portion of monthly API revenue is fed directly back into the contributor prize pool. Because the pool size determines how hard contributors compete, revenue has a compounding effect on the Product's core technology:

- API revenue → larger prize pool → stronger contributors → better kernels → a more valuable API → more revenue.
- Retained margin funds the Product team, standardized benchmark hardware, and enterprise integration.
- As the champion kernel improves, the same API becomes more valuable to existing customers at no added delivery cost — expanding margin over time.

Revenue is not a side benefit — it is the primary engine by which the Product converts commercial success into technical dominance.

10.3 Summary of Value Flows

Flow	Direction	Mechanism
API subscription fees	Customer → Product team	Direct B2B contracts
Prize pool funding	Revenue → contributor pool	Reinvested share of monthly revenue
Stronger competition	Larger pool → better kernels	Winner-take-all incentive
Champion publication	Product → public library	Open-source GitHub release each cycle
Winner payout	Prize pool → top-ranked contributor	Automated settlement, winner-take-all
Operating margin	Revenue → Product team	Retained share funds ops, hardware, R&D

11. Why This Succeeds Where Others Have Failed

Challenge	Common Failure Mode	Product Solution
LLM-as-a-judge bias	Systems that use LLMs to score other LLMs inherit the same hallucinations and biases they aim to fix.	Pure mathematical scoring: MSE, VRAM, latency on standardized hardware. No opinion involved.
One-time solution kills competition	A single breakthrough ends the race; the effort stagnates or dies.	Infinite search space: accuracy thresholds tighten, sequence lengths grow, hardware targets shift. The winner-take-all pressure never relaxes.
Too hard for new contributors	Advanced programs require PhD-level contributions from day one, limiting the talent pool.	Phase 1 is accessible to any skilled GPU engineer using existing open-source research. Difficulty scales with demonstrated progress.
No real revenue	Many research efforts produce assets with no natural buyer, making grants or speculation the only funding source.	Long-context AI inference is a multi-billion dollar market. The API product sells itself to enterprises already spending heavily on GPU compute.
Scam-vulnerable scoring	Contributors copy each other or game a judge model to farm rewards without adding value.	Randomized inputs, mandatory open-source code, kernel re-compilation, scaling curve checks, and functional fingerprinting close all major attack vectors.
Funding-only flywheel	Efforts that rely solely on external funding are vulnerable to market cycles and grant fatigue.	The reinvestment loop means the Product funds its own R&D from real customer revenue, independent of outside capital.

12. Conclusion

The quadratic scaling of self-attention is the defining bottleneck of modern AI. It is not a niche academic problem — it is the reason long-horizon AI agents, deep document understanding, and efficient enterprise inference remain economically constrained. Solving it, even partially, would unlock capabilities that enterprises are actively willing to pay for today.

Product proposes to solve this not through a single research paper, but through a continuously evolving, winner-take-all competition among the world's best GPU engineers, with all work published in the open. By

anchoring evaluation in deterministic hardware benchmarks, enforcing open-source transparency, and funding a direct reinvestment loop from real revenue, the Product avoids the failure modes — stagnant incentives, LLM judge bias, and funding-dependent economics — that have limited previous efforts.

The result is a virtuous cycle: better kernels attract more enterprise customers, more revenue grows the prize pool, a larger pool attracts more contributors, and more contributors produce even better kernels.

Product: Open-source kernels. Winner-take-all competition. Real revenue. A self-funding flywheel.

This document is a working draft whitepaper. Technical parameters (MSE thresholds, sequence lengths, scoring weights, revenue-reinvestment ratios) are subject to revision based on testnet results and community feedback. Nothing herein constitutes financial advice.